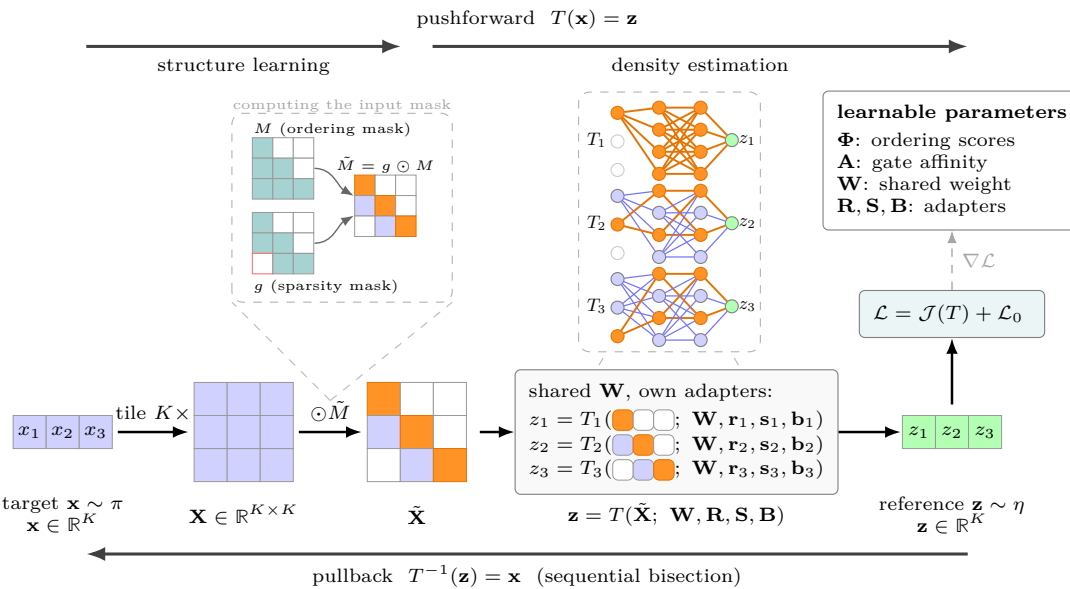

Figure 1: Overview of SSTM. The forward map pushes $\mathbf{x} \sim \pi$ to $\mathbf{z} \sim \eta$, while the inverse uses sequential bisection. The learned ordering and sparsity gates form the sparse triangular mask $\tilde{M} = g \odot M$, which is applied to the input. Each of the K masked rows is input to one map component T_k. The K components share weights $\mathbf{W}$ and use rank-one adapters $(\mathbf{r}_k, \mathbf{s}_k, \mathbf{b}_k)$. The objective $\mathcal{L} = \mathcal{J}(T) + \mathcal{L}_0$ jointly optimizes the map, ordering, and sparsity. *Cell colours:* orange, diagonal variable; blue, conditioning variable; white, masked off.